

▼ Original data

import libraries

```
from keras.layers import Input, Lambda, Dense, Flatten
from keras.models import Model
from keras.applications.vgg16 import VGG16
from keras.applications.vgg16 import preprocess_input
from keras.preprocessing import image
from keras.preprocessing.image import ImageDataGenerator
from keras.models import Sequential
import numpy as np
from glob import glob
import matplotlib.pyplot as plt

import warnings
warnings.filterwarnings("ignore", category=FutureWarning)
```

import keras

IMAGE_SIZE = [224, 224]

```
from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
!ls '/content/drive'
```

MyDrive

dataset path

```
#Give dataset path
train_path = "/content/drive/MyDrive/vgg16/men_women/train"
test_path = "/content/drive/MyDrive/vgg16/men_women/test"
```

```
from PIL import Image
import os
from IPython.display import display
from IPython.display import Image as _Imgdis
# creating a object
```

```
folder = train_path+'/men500'
```

```
onlybenignfiles = [f for f in os.listdir(folder) if os.path.isfile(os.path.join(folder, f))]
print("Working with {0} images".format(len(onlybenignfiles)))
print("Image examples: ")
```

```
for i in range(10,20):
    print(onlybenignfiles[i])
    display(_Imgdis(filename=folder + "/" + onlybenignfiles[i], width=240, height=240))
```

Working with 500 images
Image examples:
00000800.jpg



00000563.jpg



00000622.jpg



```
vgg = VGG16(input_shape=IMAGE_SIZE + [3], weights='imagenet', include_top=False)
```

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/vgg16/vgg16_weights_tf_dim_ordering_tf_kernels_notop.58889256/58889256 [=====] - 3s 0us/step



```
vgg.input
```

<KerasTensor: shape=(None, 224, 224, 3) dtype=float32 (created by layer 'input_1')>



```
for layer in vgg.layers:  
    layer.trainable = False
```

```
x = Flatten()(vgg.output)  
num_classes = 2  
prediction = Dense(num_classes, activation='softmax')(x)  
model = Model(inputs=vgg.input, outputs=prediction)  
model.summary()
```

Model: "model"

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	[(None, 224, 224, 3)]	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147584

block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590080
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1180160
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2359808
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2359808
block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0
block5_conv1 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv2 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv3 (Conv2D)	(None, 14, 14, 512)	2359808
block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0
flatten (Flatten)	(None, 25088)	0
dense (Dense)	(None, 2)	50178

```

=====
Total params: 14,764,866
Trainable params: 50,178
Non-trainable params: 14,714,688

```

```
from keras import optimizers
```

```
adam = optimizers.Adam()
model.compile(loss='binary_crossentropy',
              optimizer=adam,
              metrics=['accuracy'])
```

```
train_datagen = ImageDataGenerator(
    preprocessing_function=preprocess_input,
)
```

```
test_datagen = ImageDataGenerator(
    preprocessing_function=preprocess_input,
)
```

```
train_set = train_datagen.flow_from_directory(train_path,
                                              target_size = (224, 224),
                                              batch_size = 32,
                                              class_mode = 'categorical')
```

```
Found 995 images belonging to 2 classes.
```

```
test_set = test_datagen.flow_from_directory(test_path,
                                           target_size = (224, 224),
                                           batch_size = 32,
                                           class_mode = 'categorical')
```

```
Found 259 images belonging to 2 classes.
```

```
training the model
```

```
from keras.callbacks import ModelCheckpoint
```

```
checkpoint = ModelCheckpoint(filepath='/content/drive/MyDrive/vgg16_evaluation/mymodel.h5',
                             verbose=2, save_best_only=True)
```

```
callbacks = [checkpoint]
```

```
start = datetime.now()
```

```
model_history=model.fit(
    train_set,
    validation_data=test_set,
    epochs=10,
    # steps_per_epoch=30,
    # validation_steps=32,
    callbacks=callbacks ,verbose=2)
```

Epoch 1/10

Epoch 1: val_loss improved from inf to 1.43363, saving model to /content/drive/MyDrive/vgg16_evaluation/mymodel.h5
32/32 - 30s - loss: 0.1683 - accuracy: 0.9980 - val_loss: 1.4336 - val_accuracy: 0.9151 - 30s/epoch - 949ms/step
Epoch 2/10

Epoch 2: val_loss improved from 1.43363 to 1.24293, saving model to /content/drive/MyDrive/vgg16_evaluation/mymodel.h5
32/32 - 27s - loss: 0.0784 - accuracy: 0.9990 - val_loss: 1.2429 - val_accuracy: 0.9305 - 27s/epoch - 852ms/step
Epoch 3/10

Epoch 3: val_loss improved from 1.24293 to 1.21557, saving model to /content/drive/MyDrive/vgg16_evaluation/mymodel.h5
32/32 - 36s - loss: 0.1590 - accuracy: 0.9980 - val_loss: 1.2156 - val_accuracy: 0.9305 - 36s/epoch - 1s/step
Epoch 4/10

Epoch 4: val_loss did not improve from 1.21557
32/32 - 27s - loss: 0.1189 - accuracy: 0.9990 - val_loss: 1.4634 - val_accuracy: 0.9151 - 27s/epoch - 846ms/step
Epoch 5/10

Epoch 5: val_loss improved from 1.21557 to 1.08849, saving model to /content/drive/MyDrive/vgg16_evaluation/mymodel.h5
32/32 - 29s - loss: 0.1194 - accuracy: 0.9990 - val_loss: 1.0885 - val_accuracy: 0.9421 - 29s/epoch - 916ms/step
Epoch 6/10

Epoch 6: val_loss did not improve from 1.08849
32/32 - 26s - loss: 0.1514 - accuracy: 0.9990 - val_loss: 1.1731 - val_accuracy: 0.9498 - 26s/epoch - 822ms/step
Epoch 7/10

Epoch 7: val_loss did not improve from 1.08849
32/32 - 25s - loss: 0.0861 - accuracy: 0.9970 - val_loss: 1.4554 - val_accuracy: 0.9228 - 25s/epoch - 772ms/step
Epoch 8/10

Epoch 8: val_loss did not improve from 1.08849
32/32 - 26s - loss: 0.1951 - accuracy: 0.9980 - val_loss: 1.7916 - val_accuracy: 0.9112 - 26s/epoch - 810ms/step
Epoch 9/10

Epoch 9: val_loss improved from 1.08849 to 1.00591, saving model to /content/drive/MyDrive/vgg16_evaluation/mymodel.h5
32/32 - 26s - loss: 0.1616 - accuracy: 0.9990 - val_loss: 1.0059 - val_accuracy: 0.9498 - 26s/epoch - 806ms/step
Epoch 10/10

Epoch 10: val_loss did not improve from 1.00591
32/32 - 25s - loss: 0.1018 - accuracy: 0.9990 - val_loss: 1.1080 - val_accuracy: 0.9459 - 25s/epoch - 766ms/step
Training completed in time: 0:05:29.508975

```
model_history.history
```

```
{'loss': [0.1682899296283722,
0.07838954031467438,
0.15904904901981354,
0.11891869455575943,
0.11940443515777588,
0.15138235688209534,
0.08607212454080582,
0.19511447846889496,
0.1615639328956604,
0.10182344913482666],
'accuracy': [0.9979899525642395,
0.9989949464797974,
0.9979899525642395,
0.9989949464797974,
0.9989949464797974,
0.9989949464797974,
0.9969848990440369,
0.9979899525642395,
0.9989949464797974,
0.9989949464797974],
'val_loss': [1.4336315393447876,
1.2429298162460327,
```

```


1.215573787689209,
1.4633790254592896,
1.088494062423706,
1.1730585098266602,
1.4554328918457031,
1.7915524244308472,
1.0059056282043457,
1.1079607009887695],
'val_accuracy': [0.915057897567749,
0.9305019378662109,
0.9305019378662109,
0.915057897567749,
0.9420849680900574,
0.9498069286346436,
0.92277991771698,
0.9111968874931335,
0.9498069286346436,
0.9459459185600281]]}

```

```

_# Plot training & validation loss values
plt.plot(model_history.history['loss'])
plt.plot(model_history.history['val_loss'])
plt.title('CNN Model loss values')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc='upper right')
plt.show()

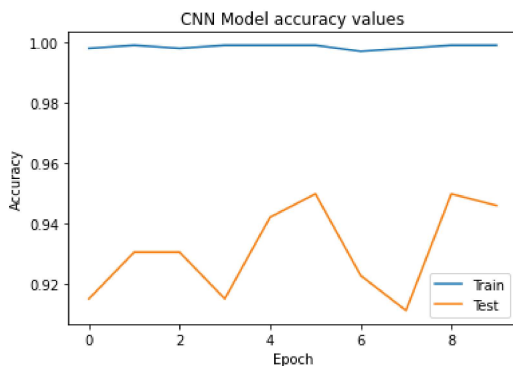
```



```

_# Plot training & validation accuracy values
plt.plot(model_history.history['accuracy'])
plt.plot(model_history.history['val_accuracy'])
plt.title('CNN Model accuracy values')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc='lower right')
plt.show()

```




Testing the model

```

img = plt.imread("/content/drive/MyDrive/vgg16/men_women/test/men130/00000001.jpg")
plt.imshow(img)

```

```
<matplotlib.image.AxesImage at 0x7f60acbedfa0>
```

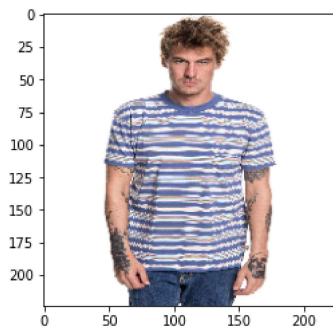


```
import tensorflow as tf
from keras.preprocessing import image
```

```
img = tf.keras.utils.load_img("/content/drive/MyDrive/vgg16/men_women/test/men130/00000001.jpg",target_size=(224,224))
img = np.asarray(img)
plt.imshow(img)
img = np.expand_dims(img, axis=0)
from keras.models import load_model
saved_model = load_model('/content/drive/MyDrive/vgg16_evaluation/mymodel.h5')
```

```
output = saved_model.predict(img)
if output[0][0] > output[0][1]:
    print("men")
else:
    print('women')
```

```
1/1 [=====] - 1s 814ms/step
men
```



```
def predict_img(img):
    img = tf.keras.utils.load_img(img,target_size=(224,224))
    img = np.asarray(img)
    plt.imshow(img)
    img = np.expand_dims(img, axis=0)
    from keras.models import load_model
    saved_model = load_model('/content/drive/MyDrive/vgg16_evaluation/mymodel.h5')

    output = saved_model.predict(img)
    if output[0][0] > output[0][1]:
        print("men")
    else:
        print('women')
```

```
predict_img("/content/drive/MyDrive/vgg16/men_women/test1/test1/women/00001315.jpg")
```

```
WARNING:tensorflow:6 out of the last 6 calls to <function Model.make_predict_func
1/1 [=====] - 0s 116ms/step
```


```
women
```



```
predict_img("/content/drive/MyDrive/vgg16/men_women/test1/test1/men/00000408.jpg")
```

```
1/1 [=====] - 0s 113ms/step
```

```
men
```



▼ For Augmented Dataset

The data augmentation is done with different code and the augmented images are saved differently (apart from original dataset). So in the following code we are considering only augmented images not original images. (Augmented code is shared with this file)

```
import librarires
```

```
import keras
from keras.layers import Input, Lambda, Dense, Flatten
from keras.models import Model
from keras.applications.vgg16 import VGG16
from keras.applications.vgg16 import preprocess_input
from keras.preprocessing import image
from keras.preprocessing.image import ImageDataGenerator
from keras.models import Sequential
import numpy as np
from glob import glob
import matplotlib.pyplot as plt
```

```
import warnings
warnings.filterwarnings("ignore", category=FutureWarning)
```

```
IMAGE_SIZE = [224, 224]
```

```
dataset path(for augmented images)
```

```
train_path = "/content/drive/MyDrive/vgg16/men_women/Augmented_dataset/train"
test_path = "/content/drive/MyDrive/vgg16/men_women/Augmented_dataset/test"
```

```
Augmented image examples
```

```
from PIL import Image
import os
from IPython.display import display
from IPython.display import Image as _Imgdis
# creating a object
```

```
folder = train_path+'/men'
```

```
onlybenignfiles = [f for f in os.listdir(folder) if os.path.isfile(os.path.join(folder, f))]
print("Working with {0} images".format(len(onlybenignfiles)))
print("Image examples: ")
```

```

for i in range(10,20):
    print(onlybenignfiles[i])
    display(_Imgdis(filename=folder + "/" + onlybenignfiles[i], width=240, height=240))

```

```

vgg = VGG16(input_shape=IMAGE_SIZE + [3], weights='imagenet', include_top=False)

```

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/vgg16/vgg16_weights_tf_dim_ordering_tf_kernels_notop.58889256/58889256 [=====] - 3s 0us/step



```

vgg.input

```

```

<KerasTensor: shape=(None, 224, 224, 3) dtype=float32 (created by layer 'input_1')>

```

```

for layer in vgg.layers:
    layer.trainable = False

```

```

x = Flatten()(vgg.output)
num_classes = 2
prediction = Dense(num_classes, activation='softmax')(x)

```



```
model = Model(inputs=vgg.input, outputs=prediction)
model.summary()
```

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 224, 224, 3)]	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147584
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590080
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1180160
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2359808
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2359808
block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0
block5_conv1 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv2 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv3 (Conv2D)	(None, 14, 14, 512)	2359808
block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0
flatten (Flatten)	(None, 25088)	0
dense (Dense)	(None, 2)	50178
Total params: 14,764,866		
Trainable params: 50,178		
Non-trainable params: 14,714,688		

```
from keras import optimizers

adam = optimizers.Adam()
model.compile(loss='binary_crossentropy',
              optimizer=adam,
              metrics=['accuracy'])

train_datagen = ImageDataGenerator(
    preprocessing_function=preprocess_input,
)

test_datagen = ImageDataGenerator(
    preprocessing_function=preprocess_input,
)

train_set = train_datagen.flow_from_directory(train_path,
                                              target_size = (224, 224),
                                              batch_size = 32,
                                              class_mode = 'categorical')

Found 949 images belonging to 2 classes.

test_set = test_datagen.flow_from_directory(test_path,
                                           target_size = (224, 224),
```

```
batch_size = 32,
class_mode = 'categorical')
```

```
Found 260 images belonging to 2 classes.
```

```
from keras.callbacks import ModelCheckpoint
```

```
checkpoint = ModelCheckpoint(filepath='/content/drive/MyDrive/vgg16_evaluation/vgg16_Augmented.h5',
                             verbose=2, save_best_only=True)
```

```
callbacks = [checkpoint]
```

```
start = datetime.now()
```

```
model_history=model.fit(
    train_set,
    validation_data=test_set,
    epochs=10,
    # steps_per_epoch=30,
    # validation_steps=32,
    callbacks=callbacks ,verbose=2)
```

```
Epoch 1/10
```

```
Epoch 1: val_loss improved from inf to 3.15409, saving model to /content/drive/MyDrive/vgg16_evaluation/vgg16_Augmented.h5
30/30 - 18s - loss: 0.5643 - accuracy: 0.9579 - val_loss: 3.1541 - val_accuracy: 0.7692 - 18s/epoch - 590ms/step
Epoch 2/10
```

```
Epoch 2: val_loss improved from 3.15409 to 3.10146, saving model to /content/drive/MyDrive/vgg16_evaluation/vgg16_Augmented.h5
30/30 - 18s - loss: 0.1411 - accuracy: 0.9874 - val_loss: 3.1015 - val_accuracy: 0.7692 - 18s/epoch - 603ms/step
Epoch 3/10
```

```
Epoch 3: val_loss did not improve from 3.10146
30/30 - 20s - loss: 0.0230 - accuracy: 0.9989 - val_loss: 3.3371 - val_accuracy: 0.7654 - 20s/epoch - 679ms/step
Epoch 4/10
```

```
Epoch 4: val_loss did not improve from 3.10146
30/30 - 16s - loss: 0.0154 - accuracy: 1.0000 - val_loss: 3.2833 - val_accuracy: 0.7769 - 16s/epoch - 545ms/step
Epoch 5/10
```

```
Epoch 5: val_loss did not improve from 3.10146
30/30 - 16s - loss: 0.0014 - accuracy: 1.0000 - val_loss: 3.2312 - val_accuracy: 0.7692 - 16s/epoch - 543ms/step
Epoch 6/10
```

```
Epoch 6: val_loss did not improve from 3.10146
30/30 - 16s - loss: 0.0011 - accuracy: 1.0000 - val_loss: 3.2112 - val_accuracy: 0.7769 - 16s/epoch - 536ms/step
Epoch 7/10
```

```
Epoch 7: val_loss did not improve from 3.10146
30/30 - 16s - loss: 6.0736e-05 - accuracy: 1.0000 - val_loss: 3.2032 - val_accuracy: 0.7731 - 16s/epoch - 537ms/step
Epoch 8/10
```

```
Epoch 8: val_loss did not improve from 3.10146
30/30 - 16s - loss: 4.8455e-05 - accuracy: 1.0000 - val_loss: 3.2020 - val_accuracy: 0.7731 - 16s/epoch - 535ms/step
Epoch 9/10
```

```
Epoch 9: val_loss did not improve from 3.10146
30/30 - 16s - loss: 4.1059e-05 - accuracy: 1.0000 - val_loss: 3.2013 - val_accuracy: 0.7731 - 16s/epoch - 538ms/step
Epoch 10/10
```

```
Epoch 10: val_loss did not improve from 3.10146
30/30 - 16s - loss: 3.6654e-05 - accuracy: 1.0000 - val_loss: 3.2008 - val_accuracy: 0.7731 - 16s/epoch - 534ms/step
```

```
model_history.history
```

```
{'loss': [0.564281165599823,
0.14114993810653687,
0.022950368002057076,
0.015424290671944618,
0.001445319503545761,
0.001145491492934525,
6.073633267078549e-05,
4.845533112529665e-05,
4.105871994397603e-05,
```

```


3.6654080759035423e-05],
'accuracy': [0.9578503966331482,
0.98735511302948,
0.9989462494850159,
1.0,
1.0,
1.0,
1.0,
1.0,
1.0,
1.0],
'val_loss': [3.154093027114868,
3.1014645099639893,
3.3371238708496094,
3.2832515239715576,
3.231226682662964,
3.211181402206421,
3.203233480453491,
3.2019853591918945,
3.2012991905212402,
3.200820207595825],
'val_accuracy': [0.7692307829856873,
0.7692307829856873,
0.7653846144676208,
0.7769230604171753,
0.7692307829856873,
0.7769230604171753,
0.7730769515037537,
0.7730769515037537,
0.7730769515037537,
0.7730769515037537]]}

```

```

_# Plot training & validation loss values
plt.plot(model_history.history['loss'])
plt.plot(model_history.history['val_loss'])
plt.title('CNN Model loss values')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc='upper right')
plt.show()


```



```

_# Plot training & validation accuracy values
plt.plot(model_history.history['accuracy'])
plt.plot(model_history.history['val_accuracy'])
plt.title('CNN Model accuracy values')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc='lower right')
plt.show()

```



```
import tensorflow as tf
from keras.preprocessing import image

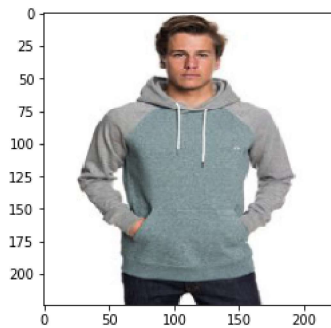
def predict_img(img):
    img = tf.keras.utils.load_img(img, target_size=(224, 224))
    img = np.asarray(img)
    plt.imshow(img)
    img = np.expand_dims(img, axis=0)
    from keras.models import load_model
    saved_model = load_model('/content/drive/MyDrive/vgg16_evaluation/vgg16_Augmented.h5')

    output = saved_model.predict(img)
    if output[0][0] > output[0][1]:
        print("men")
    else:
        print('women')
```

```
predict_img("/content/drive/MyDrive/vgg16/men_women/test1/test1/men/00000399.jpg")
```

```
1/1 [=====] - 0s 120ms/step
```

```
men
```



```
predict_img("/content/drive/MyDrive/vgg16/men_women/test1/test1/women/00001318.jpg")
```

```
1/1 [=====] - 0s 123ms/step
```

```
women
```

